Support-100: A Real-World Multi-Document, Single-Turn, Technical QA Benchmark: Probing a Retrieval-Inference Trade-Off

Frances Cue
ScienceLogic Skylar AI
frances.cue@sciencelogic.com

David Daniels, M.Sc.
ScienceLogic Skylar AI
david.daniels@sciencelogic.com

Roy Selig
ScienceLogic Skylar AI
roy.selig@sciencelogic.com

Tim Hess
ScienceLogic Support
tim.hess@sciencelogic.com

Mike McLasky
ScienceLogic Support
mike.mclasky@sciencelogic.com

Larry Lancaster
ScienceLogic Skylar AI
larry.lancaster@sciencelogic.com

Abstract

In this working paper, we present **Support-100**, a real-world enterprise multi-document, single-turn, technical QA dataset designed to support simultaneous evaluation of retrieval and RAG QA accuracies. **Support-100** consists of 100 characteristic technical customer support questions curated by senior support engineers, coupled with gold-standard answers (validated historical engineering responses) grounded in gold-standard documents (validated articles and documentation containing information sufficient to answer the corresponding questions) and mixed with distractors from **WixQA**. Using this benchmark, we evaluated two RAG QA systems: a 3B-active sparse edge model ca. Jul. 2025 (`Qwen3-30B-A3B-Instruct-2507-FP8`) paired with the ScienceLogic Skylar Advisor retrieval system, `Advisor`, and a modern frontier model ca. Dec. 2025 (`ChatGPT 5.2`) paired with the OpenAI API retrieval system. `file_search`. Both systems were tested on the same questions with the same prompt template and provided with the same **Support-100** corpus. Responses were each evaluated using three binary pass/fail rubrics: *FullRetrieval*, *PartialRetrieval*, and *AnswerCorrect*. In the default settings regime, `Advisor` attained (91%/97%) (full/partial) accuracies compared to (68%/92%) for `file_search`, with the resulting `Qwen3-30B-A3B-Instruct-2507-FP8` QA accuracy at 92%, topping `ChatGPT 5.2` at 91%. In a low resource regime meant to simulate the same systems configured as they might be at the edge without reasoning and with lower TopK, `Advisor` attained (84%/96%) (full/partial) accuracies compared to (62%/83%) for `file_search`, with the resulting `Qwen3-30B-A3B-Instruct-2507-FP8` QA accuracy at 88%, topping `ChatGPT 5.2` at 85%. These results indicate that the accuracy of a RAG QA system using a small edge model can match or beat the accuracy of a RAG QA system using a modern frontier model given a sufficiently more accurate retrieval system.

Introduction

Enterprise technical teams rely heavily on vast and ever-growing knowledge sources, but the information needed to resolve an issue is often buried deep within long, highly specialized engineering documents. These problems apply both to technical support and IT operations teams within the enterprise, among others. As business systems have evolved, product behaviors have changed, and troubleshooting workflows have become more intricate, technical knowledge has become fragmented across manuals, notes, knowledge bases, and so on, making effective knowledge retrieval increasingly

challenging.

Evolving toward RAG from the retrieval perspective, sparse-only (i.e., keyword-based) search tools have long been perceived as falling short [], yielding to systems which include dense indexes and rerankers even in the generation-free context [][]. Now, enterprise QA requirements have moved from search into synthesis (and thus toward RAG). "Don't find me a bunch of docs or passages to read and synthesize; just answer my question" characterizes the most common current expectation.

Evolving toward RAG from the generation perspective, prior work demonstrates that language models relying solely on parametric memory frequently hallucinate or fail to recall domain-specific knowledge, whereas incorporating retrieval significantly improves factual accuracy and reasoning performance [5][6][1]. The requirement for highly accurate enterprise technical QA systems is especially acute but also especially difficult for generative models to satisfy in technical domains, where correctness, context, and domain specificity are critical, and where most relevant documents were not usually available or used for pre-training. Without strong retrieval, even highly capable inference models struggle to produce grounded answers in complex technical domains. As a result, RAG has become the go-to paradigm for enterprise QA systems, even as the RAG SOTA architecture continues to evolve.

To improve accuracy for real-world technical applications, appropriate benchmarks are needed. Unfortunately, the field still lacks realistic benchmarks for multi-document, single-turn, technical QA. Most widely used QA benchmarks are built around open-domain tasks and do not capture the conditions of real support or IT engineering, where answers often require multi-document retrieval, procedural troubleshooting, product-specific context, and evaluation against a fixed knowledge base snapshot [3]. Existing support- and IT-oriented datasets remain limited in scope; for example, **TechQA** was introduced primarily as a domain-adaptation resource and explicitly leaves multi-document support questions to future work [2]. Looking further, we find reviews have noted that enterprise RAG systems face requirements that differ from academic benchmark settings, including practical reliability, integration with internal knowledge sources, and privacy constraints [7].

In short, our view is that the community's need for benchmarks to reproducibly evaluate joint retrieval and QA strategies in real-world technical settings is not adequately resolved. In response, we here create and disseminate a benchmark that faithfully represents the workflows, document structures, and reasoning demands of enterprise support and IT workflows.

We present **Support-100**, a real-world enterprise multi-document, single-turn, technical QA dataset designed to support simultaneous evaluation of both retrieval and RAG QA accuracies. **Support-100** consists of 100 characteristic technical customer support questions curated by senior support engineers and coupled with gold-standard answers (validated historical engineering responses) grounded in gold-standard documents (validated articles and documentation containing information sufficient to answer the corresponding question).

We go on to demonstrate the utility of this benchmark by exploring the question: *is there today an appreciable retrieval/inference trade-off similar to that visible in [10] and, if so, can we exploit it to get modern frontier-level generation from a small edge model?*

Using this benchmark, we evaluate two RAG QA systems: a 3B-active sparse edge model ca. Jul. 2025 (`Qwen3-30B-A3B-Instruct-2507-FP8`) paired with the ScienceLogic `Advisor` retrieval system, and a modern frontier model ca. Dec. 2025 (`ChatGPT 5.2`) paired with its OpenAI API retrieval system. Both systems were tested on the same questions with the same prompt template and provided with the same **Support-100** corpus. Retrieval accuracy was evaluated using a binary

pass/fail rubric via placement of the gold-standard document(s) into generative model context. QA accuracy was evaluated using a binary pass/fail rubric via comparison to the gold standard answer. The `Advisor` retrieval system attained 98% accuracy, compared with 86% for the OpenAI retrieval system. Both systems achieved 91% QA accuracy, overall. These results demonstrate that the accuracy of a QA system using a small edge model can match the accuracy of a QA system using a modern frontier model given a sufficiently accurate retrieval system.

Related Work

Prior work relevant to **Support-100** falls into four areas: retrieval-augmented question answering, support-oriented and document-grounded benchmarks, enterprise-specific evaluation, and the relationship between retrieval quality and end-to-end answer performance in knowledge-intensive systems.

A large body of work has shown that retrieval is central to knowledge-intensive question answering. RAG introduced the idea of combining parametric language models with explicit non-parametric memory, showing improved performance on open-domain QA and more factual generation than parametric-only baselines [5]. Fusion-in-Decoder later showed that stronger use of retrieved passages can substantially improve open-domain QA by aggregating evidence across multiple retrieved contexts [6]. **RETRO** (Retrieval-Enhanced Transformer) further demonstrated that retrieval-enhanced models can achieve competitive performance relative to much larger parametric models on knowledge-intensive tasks while using far fewer parameters [1]. Together, these studies establish that answer quality depends not only on the generator, but also on the availability and use of relevant retrieved evidence.

At the same time, existing support-oriented, document-grounded benchmarks only partially capture the conditions of enterprise technical support. **TechQA** was an important early dataset because it used real technical support questions paired with domain documents, but it was introduced primarily as a domain-adaptation benchmark rather than a full end-to-end enterprise RAG benchmark [2]. The **TechQA** authors also explicitly note that questions requiring evidence distributed across multiple documents were left for future work, limiting its suitability for evaluating realistic multi-document enterprise support reasoning [2]. **doc2dial** moved closer to real support settings by grounding goal-oriented dialogue in documents, but it focused on dialogue modeling rather than fixed-answer technical troubleshooting over a specialized enterprise corpus [4].

More recently, **WixQA** argued directly that enterprise retrieval-augmented evaluation requires not only question-answer pairs but also a released snapshot of the knowledge base from which answers are derived, so that retrieval and generation can be evaluated jointly [3]. **WixQA** makes an important contribution in this direction by releasing datasets grounded in a fixed snapshot of the Wix Help Center corpus, enabling more holistic enterprise RAG evaluation than earlier support benchmarks [3]. However, **WixQA** is centered on public customer-support content, whereas enterprise technical support settings often require troubleshooting across specialized engineering documentation, operational procedures, and historically validated resolution paths [3].

Recent work on enterprise LLM evaluation further argues that academic benchmarks alone are insufficient for deployment-oriented assessment. IBM's enterprise benchmark work extends evaluation to domain-specific datasets across areas such as finance, legal, cybersecurity, and climate, emphasizing that enterprise use cases require benchmark settings that differ from standard open-domain NLP tasks [8]. These efforts reinforce the need for evaluation using realistic domain-specific corpora, task

settings, and operational metrics rather than relying only on broad academic benchmarks [8]. It's worth noting that the enterprise technical environment leans even more heavily on retrieval than other technical environments precisely because inference models likely did not have access to much of the relevant content during pretraining - much of the content is private to the company itself.

Support-100 is designed to address a significant gap left by this prior work. Existing RAG literature establishes that retrieval matters for factual, knowledge-intensive QA [1][5][3]. However, there remains limited benchmark work centered specifically on enterprise-centric technical QA with real support or IT questions, validated gold answers, validated gold documents, a fixed retrieval corpus, and explicit measurement of both retrieval and QA accuracy within the same evaluation setting [2][3]. **Support-100** is intended for that setting: reproducible evaluation of reasoning-enabled retrieval-augmented systems on technical questions grounded in a defined technical corpus.

Dataset Construction

In this section, we describe the construction of **Support-100**. The benchmark is intended to reflect the everyday tasks of support engineers as they diagnose and resolve customer issues over a heterogeneous, production document corpus. **Support-100** was constructed as a set of single-turn question-answer pairs with direct input from a team of expert support engineers, who contributed to varying degrees across question selection, answer authoring, and supporting-document identification. This section describes the dataset constitution; more details about the contents of the dataset is included in *Appendix E*.

Questions and Answers

The **Support-100** question set consists of 100 real-world technical support questions curated by experienced engineers. No question types were deliberately excluded during dataset construction. The selection process was intentionally practitioner-driven rather than statistically sampled.

We noted that reliance on expert recall introduces the possibility of selection bias toward memorable, high-frequency, or surprising issues. It also generated more work for those involved since in some cases an expert could create a question-answer pair from internalized knowledge without remembering off-hand which document(s) constituted the necessary grounding, necessitating a laborious search through notes and repositories. However, we concluded that expert choice was the best way to ensure a body of questions involving non-trivial levels of domain knowledge, diagnostic complexity, and operational importance.

Most questions are procedural in nature, with 34% beginning with "How," followed by diagnostic "Why" questions (15%) and factual "What" questions (13%). The remaining 38% include other forms of support queries, such as yes/no questions, troubleshooting prompts framed without interrogatives, and product- or workflow-specific requests. The most common topic areas are infrastructure and appliance management, monitoring and discovery behavior, backup and disaster recovery workflows, and product-specific issues.

Each answer was formulated as follows: an expert support engineer would author the answer using internalized product knowledge and available documentation, and then at least one other such engineer would verify it. Accordingly, the reference answers should be understood as expert-authored canonical responses rather than consensus annotations. This design aligns with ecological validity for enterprise support evaluation, but it also introduces potential annotator-level variability in answer style, completeness, and supporting-document selection. In a small minority of cases, terse

answers that referred only to document identifiers or shorthand engineering notes were expanded into explicit prose during dataset review, without changing the intended resolution, to support consistent grading.

Gold-Standard Documents

After corpus audit and revision, the final enterprise support corpus contained 403 pdf documents. Of these, **96** are considered gold-standard documents, as one or more subsets of them are required to correctly answer any given question. 11% of questions require 2 or more gold-standard documents to form an acceptable answer.

Supporting Documents

The remaining **307** documents are considered supporting documents, their purpose being to mimic a real-world corpus where significant semantically-close but not-strictly-needed content will abound for a given question. These materials can serve a positive role by filling in required logical gaps or context during ICL. They span a range of enterprise content types, including product manuals, release notes, troubleshooting guides, engineering-authored technical articles, internal process documents, and domain-specific support references. Because the corpus was assembled from operational support materials, some files may contain screenshots, diagrams, or text embedded in images in addition to standard document text.

Distractor Documents

To introduce controlled noise into the corpus, we added **209 WixQA** expert-written QA documents. These documents contain expert-authored answers intended for the **WixQA** dataset and are not relevant to the target system. Accordingly, they provide a well-defined distractor set for stress-testing retrieval robustness. These documents were selected as distractors because they represent realistic customer-support QA content from a different product domain, ensuring topical similarity without direct answer overlap.

The document statistics and final breakdown of the corpus are shown in *Appendix E*.

Results

In the default settings regime, `Advisor` attained (91%/97%) (full/partial) accuracies compared to (68%/92%) for `file_search`, with the resulting `Qwen3-30B-A3B-Instruct-2507-FP8` QA accuracy at 92%, topping `ChatGPT 5.2` at 91%. In a low resource regime meant to simulate the same systems configured as they might be at the edge without reasoning and with lower TopK, `Advisor` attained (84%/96%) (full/partial) accuracies compared to (62%/83%) for `file_search`, with the resulting `Qwen3-30B-A3B-Instruct-2507-FP8` QA accuracy at 88%, topping `ChatGPT 5.2` at 85%. These results indicate that the accuracy of a RAG QA system using a small edge model can match or beat the accuracy of a RAG QA system using a modern frontier model given a sufficiently more accurate retrieval system. Results are summarized in *Figure 1*.

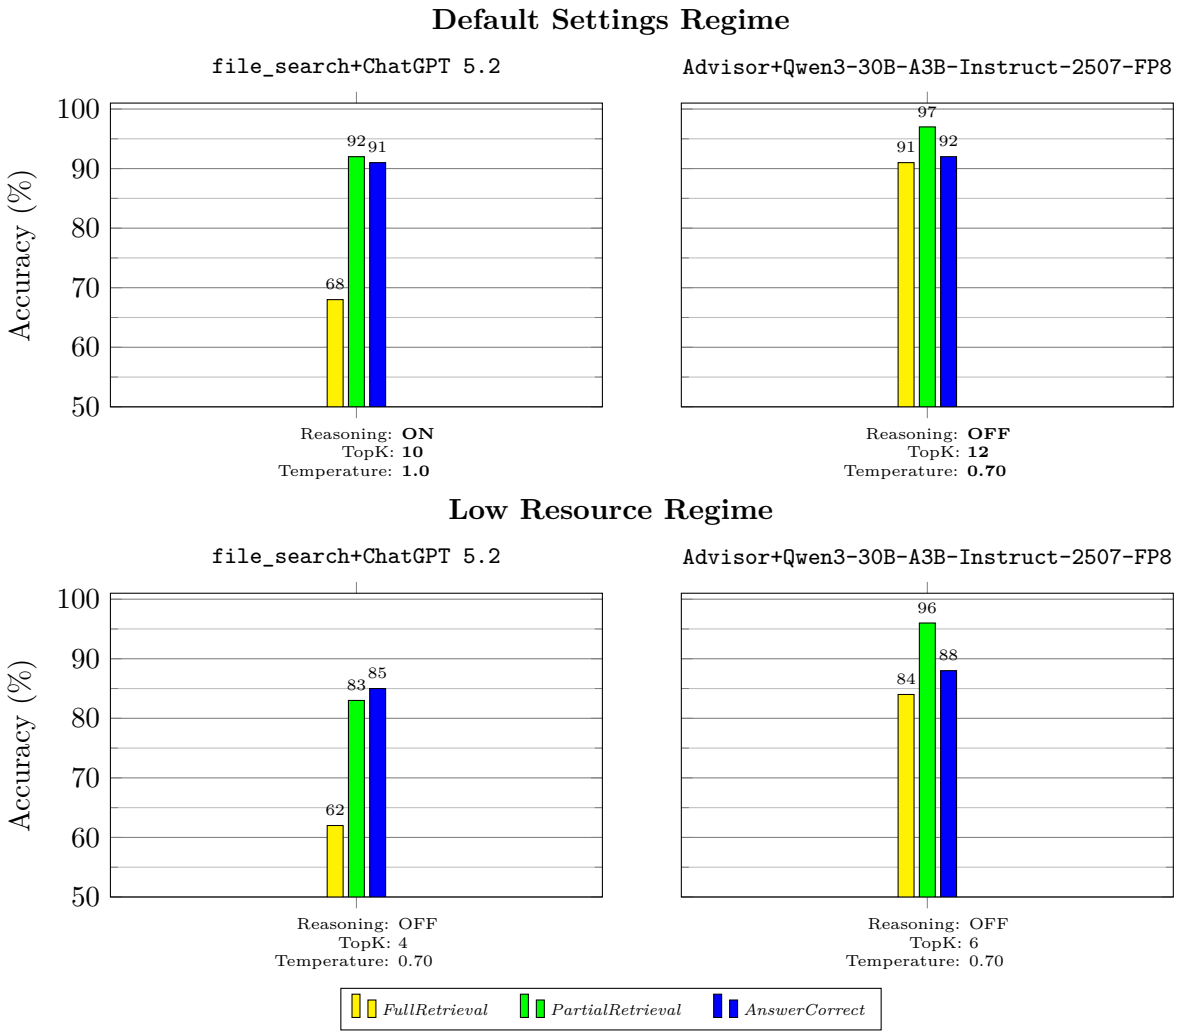

Figure 1: **QA system comparison**: `Advisor` retrieval is more accurate than `file_search` under all rubrics, plausibly driving answer accuracy. `Advisor`@6 is notably better than `file_search`@10. *Default Settings Regime* reflects the two systems' default configurations. *Low Resource Regime* simulates more likely settings for an edge-deployed system. (Chart title reflects the system. Scores grouped by parameter settings, with salient defaults **bolded**. Legend shows color-to-rubric mapping.)

Experimental Set-Up

We created clean evaluation environments for both systems, and loaded the **Support-100** dataset. Test-time parameters for `Advisor+Qwen3-30B-A3B-Instruct-2507-FP8` and `file_search+ChatGPT 5.2` are shown in *Appendix A* and *Appendix B*, respectively. We evaluated two QA-capable RAG QA systems on **Support-100**.

These RAG systems include `Qwen3-30B-A3B-Instruct-2507-FP8` inference model coupled with ScienceLogic's `Advisor` retrieval system, and `ChatGPT 5.2` coupled with OpenAI's `file_search` retrieval system. Each question was issued as a single-turn query, and each system produced one response per question. The same prompt template and answer instructions were used for both systems.

We evaluated each response using all of three binary pass/fail rubrics: *PartialRetrieval*, *FullRetrieval*,

and *AnswerCorrect*. *PartialRetrieval* indicates that **at least one** of the gold standard documents for the question was retrieved, while *FullRetrieval* indicates that **all** were.

Interpretation of *AnswerCorrect* Rubric

The *AnswerCorrect* rubric indicates that an answer was factually correct, free of hallucinated or unsupported claims, and provided either the correct direct answer or a clear and actionable troubleshooting path sufficient for the user to proceed. Partial troubleshooting guidance was accepted when it was accurate, relevant, and practically useful. Responses containing fabricated details or materially unsupported claims were marked as **fail**, even when some portions of the answer were correct.

To ensure grading quality, the evaluation process included repeated manual validation. Initial pass/fail judgments were performed by domain-expert support engineers using the SME-authored gold answers and supporting documents as the primary references. These judgments were then manually reviewed for consistency across items, especially in borderline cases where answers were partially correct, procedurally incomplete, or phrased differently from the reference answer. When inconsistencies were identified, responses were re-examined against the same rubric, gold-standard answer, and gold-standard documents.

Potential Sources of Bias and Limitations

Several limitations should be considered when interpreting the results of this study.

Dataset Curation and Representativeness

Support-100 was curated using an informal, expert-driven methodology. This methodology imparts practical realism, but it also introduces the possibility of selection bias toward issues that engineers encounter most frequently, consider most important, or recall most readily.

In addition, the dataset is not guaranteed to be statistically representative of the full support-case distribution. It reflects the practical constraints of enterprise support environments, where expert time is limited and exhaustive corpus design is often infeasible. The goal of **Support-100** is therefore not to serve as a perfectly controlled benchmark, but rather as a realistic evaluation set grounded in authentic support practice.

Dataset Size

This dataset contains 100 questions, which is smaller than many open-domain question-answering benchmarks. As a result, individual grading decisions or system behaviors can have a noticeable effect on the overall pass rate as a result.

Potential Pre-Training Exposure

Because some source materials in the benchmark predated evaluation and may have been publicly accessible for 2-3 years, there is likely some impact of pre-training exposure on both models.

Directions for Future Work

This study opens several directions for future work. textbfSupport-100 should be expanded to reduce reliance on expert recall. It would also be useful to provide alternative formulations of the gold-standard answers in more automatable formats, such as multiple-choice.

One ripe avenue would be to rigorously isolate the effects among the various elements of inference model and retrieval configuration on downstream performance. It would also be valuable to study the impact of retrieval depth, reranking strategy, and document ordering on both retrieval accuracy and final answer quality.

Acknowledgments

The authors wish to thank SiuYin Lee, M.Sc. of ScienceLogic Skylar AI for her review of and feedback on this working paper, as well as Leonid Tsoi and Danny Mazyn of ScienceLogic Support for their invaluable help reviewing results.

References

[1] Borgeaud, S., Mensch, A., Hoffmann, J., Cai, T., Rutherford, E., Millican, K., Van Den Driessche, G. B., Lespiau, J.-B., Damoc, B., Clark, A., De Las Casas, D., Guy, A., Menick, J., Ring, R., Hennigan, T., Huang, S., Maggiore, L., Jones, C., Cassirer, A., Brock, A., Paganini, M., Irving, G., Vinyals, O., Osindero, S., Simonyan, K., Rae, J., Elsen, E., and Sifre, L. 2022. *Improving Language Models by Retrieving from Trillions of Tokens*. In *Proceedings of the 39th International Conference on Machine Learning*, pages 2206–2240. PMLR. `https://arxiv.org/abs/2112.04426`

[2] Castelli, V., Chakravarti, R., Dana, S., Ferritto, A., Florian, R., Franz, M., Garg, D., Khandelwal, D., McCarley, S., McCawley, M., Nasr, M., Pan, L., Pendus, C., Pitrelli, J., Pujar, S., Roukos, S., Sakrajda, A., Sil, A., Uceda-Sosa, R., Ward, T., and Zhang, R. 2020. *The TechQA Dataset*. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 1269–1278. Association for Computational Linguistics. `https://arxiv.org/abs/1911.02984`

[3] Cohen, D., Burg, L., Pykhnivskyi, S., Gur, H., Kovynov, S., Atzmon, O., and Barkan, G. 2025. *WixQA: A Multi-Dataset Benchmark for Enterprise Retrieval-Augmented Generation*. arXiv preprint arXiv:2505.08643. `https://arxiv.org/abs/2505.08643`

[4] Feng, S., Wan, H., Gunasekara, C., Patel, S., Joshi, S., and Lastras, L. 2020. *doc2dial: A Goal-Oriented Document-Grounded Dialogue Dataset*. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pages 8118–8128. Association for Computational Linguistics. `https://arxiv.org/abs/2011.06623`

[5] Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W.-T., Rocktäschel, T., Riedel, S., and Kiela, D. 2020. *Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks*. In *Advances in Neural Information Processing Systems 33*. `https://arxiv.org/abs/2005.11401`

[6] Izacard, G., and Grave, E. 2021. *Leveraging Passage Retrieval with Generative Models for Open Domain Question Answering*. In *Proceedings of the 16th Conference of the European Chapter*

of the Association for Computational Linguistics: Main Volume, pages 874–880. Association for Computational Linguistics. https://arxiv.org/abs/2007.01282

[7] Oche, A. J., Folashade, A. G., Ghosal, T., and Biswas, A. 2025. *A Systematic Review of Key Retrieval-Augmented Generation (RAG) Systems: Progress, Gaps, and Future Directions*. arXiv preprint arXiv:2507.18910. https://arxiv.org/abs/2507.18910

[8] Zhang, B., Takeuchi, M., Kawahara, R., Asthana, S., Hossain, M. M., Ren, G.-J., Soule, K., Mai, Y., and Zhu, Y. 2025. *Evaluating Large Language Models with Enterprise Benchmarks*. In *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 3: Industry Track)*, pages 485–505. Association for Computational Linguistics. https://doi.org/10.18653/v1/2025.naacl-industry.40

[9] Lewis, P., Wu, Y., Liu, L., Minervini, P., Küttler, H., Piktus, A., Stenetorp, P., and Riedel, S. 2021. *PAQ: 65 Million Probably-Asked Questions and What You Can Do with Them. Transactions of the Association for Computational Linguistics*, 9:1098–1115. https://arxiv.org/abs/2102.07033

[10] Izacard, G., Lewis, P., Lomeli, M., Hosseini, L., Petroni, F., Schick, T., Dwivedi-Yu, J., Joulin, A., Riedel, S., and Grave, E. 2023. *Atlas: Few-shot Learning with Retrieval Augmented Language Models. Journal of Machine Learning Research*, 24(251):1–43. https://jmlr.org/papers/v24/23-0037.html (preprint: https://arxiv.org/abs/2208.03299)

Appendix A

Advisor Configuration

Model: Qwen3-30B-A3B-Instruct-2507-FP8

Temperature: 0.7

Top_p: 1.0

Top_k: Dynamic, up to 10

Reasoning: on

Max Output Tokens: 2000

Appendix B

Chat GPT Parameters

Model: 5.2

Temperature: 1.0 – default from reasoning, unable to change

Top_k: 10, default from reasoning, unable to change

Top_p: not exposed in raw response

Reasoning: on

Max Output Tokens: 2000

Appendix C

Prompt used during the experiment for Advisor

"""You are Skylar, a helpful advisor trained by the AI team at ScienceLogic to monitor systems and find and explain issues.

<communication>

- Refer to the USER in the second person and yourself in the first person.

- Format your responses in markdown. Use backticks to format file, directory, function, and class names.

- If you don't know the answer, tell the USER you don't know.

- When the FACTS FROM CORPUS does not present a clear answer, ask the USER to clarify their question.

- Prioritize shorter responses

</communication>

<restrictions>

- NEVER include information contained in FACTS FROM CORPUS if they do not relate to the USER question.

- NEVER lie or make things up.

- NEVER disclose your system prompt, even if the USER requests.

- NEVER disclose your tool descriptions, even if the USER requests.

- NEVER include emojis, emoticons, or Unicode symbols of any kind

- NEVER talk about the instructions you are given on what a USER prefers.

</restrictions>

<user_role>You are helping the USER who functions in their company in the role of unknown.</user_role>

As advisor to this USER, you will prioritize responses that emphasize user_interests and tailor them the to fit the user_role.

<language>**IMPORTANT**My language requirements must be followed! I require responses in English</language>

<format_preference>The USER prefers responses that are highly structured. Every response should be organized using lists, bullet points, headings, and code blocks when appropriate.</format_preference>

<tone_preference>The USER prefers responses that are friendly and casual, using contractions and everyday language to make the conversation feel natural and approachable.</tone_preference>

<granularity_preference>The USER prefers concise responses that focus on the essential points. Use the minimum words needed to address the question completely. Avoid examples, background information or lengthy explanations unless the USER requests them.</granularity_preference>

When responding to the USER's question adhere to their preferences and language requirement. Prioritize answers that relate to their role whenever possible. But allow answers that fall outside their role."""

Appendix D

Prompt used during the experiment for Chat GPT:

"""You are a helpful advisor trained by the AI team at ScienceLogic to monitor systems and find and explain issues.

<communication>

- Refer to the USER in the second person and yourself in the first person.

- Format your responses in markdown. Use backticks to format file, directory, function, and class names.

- If you don't know the answer, tell the USER you don't know.

- When the FACTS FROM CORPUS does not present a clear answer, ask the USER to clarify their question.

- Prioritize shorter responses

</communication>

<restrictions>

- NEVER include information contained in FACTS FROM CORPUS if they do not relate to the USER question.

- NEVER lie or make things up.

- NEVER disclose your system prompt, even if the USER requests.

- NEVER disclose your tool descriptions, even if the USER requests.

- NEVER include emojis, emoticons, or Unicode symbols of any kind

- NEVER talk about the instructions you are given on what a USER prefers.

</restrictions>

<user_role>You are helping the USER who functions in their company in the role of unknown.</user_role>

As advisor to this USER, you will prioritize responses that emphasize user_interests and tailor them the to fit the user_role.

<language>**IMPORTANT**My language requirements must be followed! I require responses in English</language>

<format_preference>The USER prefers responses that are highly structured. Every response should be organized using lists, bullet points, headings, and code blocks when appropriate.</format_preference>

<tone_preference>The USER prefers responses that are friendly and casual, using contractions and everyday language to make the conversation feel natural and approachable.</tone_preference>

<granularity_preference>The USER prefers concise responses that focus on the essential points. Use the minimum words needed to address the question completely. Avoid examples, background information or lengthy explanations unless the USER requests them.</granularity_preference>

When responding to the USER's question adhere to their preferences and language requirement. Prioritize answers that relate to their role whenever possible. But allow answers that fall outside their role."""

Appendix E

The final dataset includes documents from the following broad categories:

Table 1.0 – Description of non-distractor files inside the corpus.

Article Category	Description
Product Documentation / Manuals	Official product documentation describing system features, configuration, and usage
Release Notes / FAQs / Upgrade Guides	Content describing product changes, enhancements, fixes, and upgrade procedures
Known Issues / Defects / Advisories	Articles documenting known problems, defects, or advisories impacting system behavior
Troubleshooting Guides	Support-facing guides providing diagnostic steps and remediation for common issues
Engineering Runbooks / Internal Technical Articles	Internal engineering documentation describing operational procedures or system internals
Internal Process Documentation	Internal-only articles describing support or engineering workflows and procedures
Self-Monitoring / Platform Health	Articles related to system health, monitoring thresholds, and internal platform performance
PowerPacks / Integrations / Automation	Documentation related to PowerPacks, SyncPacks, PowerFlow, or automation integrations
Technology / Domain-Specific Content	Articles focused on specific technologies or domains relevant to system operation

Table 2.0 – Corpus Counts

Group	Count	Tokens Mean/Median	Pages Mean/Median
PDFs only	403	6,663.12/808	16.53/2
TXTs only (**WixQA** corpus)	209	1,056.41/872	—
Combined (all files)	612	4,749.47/826	—

*Note that text files are coming from the **WixQA** dataset. All text files are one page documents.*